

Vehicle Detection System – Technical Report

Model: YOLOv10 (Base)

Executive Summary

We developed a real-time vehicle detection and classification system leveraging the **YOLOv10 (Base)** model, managed through a user-friendly Streamlit GUI. The system is optimized for high-performance deployment on dual NVIDIA T4 GPUs, achieving a detection rate of **≈93 FPS** across 11 vehicle classes with high accuracy (**mAP@50 ≈0.752**). This makes it an ideal, production-ready solution for high-efficiency traffic monitoring and intelligent transportation applications.

YOLOv10 was selected for its superior balance of speed and accuracy, essential for real-time applications:

- **Real-time Performance:** Its single-stage architecture ensures low latency (approximately 7.7 ms/image on dual T4).
- **High Accuracy:** The model features improved feature extraction over its predecessor (YOLOv8), making it highly effective for challenging targets like small vehicles (motorcycle) and complex shapes (truck).
- **Efficiency:** A streamlined, optimized backbone and a reduced parameter count ($\approx 11.1M$) enable faster inference without sacrificing performance.
- **Multi-scale Detection:** The architecture effectively handles a wide range of vehicle sizes, from large buses to small motorcycles in dense traffic.

2.2 Chosen Variant Details

- **Model:** YOLOv10 (base / medium-sized variant)
- **Parameters:** $\approx 11.1M$
- **Input Size:** 640×640 (Can be configured to 416×416 for maximum speed.)
- **Architecture:** Optimized CSP backbone, enhanced PANet neck, and an anchor-free detection head.

2.3 Alternatives Considered

Model	Pros	Cons	Decision
YOLOv8	Stable, good car detection	Slower, less accurate on small vehicles	Not chosen ▾
Faster R-CNN	High accuracy (two-stage)	High latency, non-real-time	Rejected ▾
EfficientDet	Good efficiency	Complex pipeline, harder to tune	Rejected ▾

RetinaNet	Strong on rare classes	Slower than YOLO variants	Rejected ▾
-----------	------------------------	---------------------------	------------

3. Tracking Algorithm (Optional)

Vehicle tracking is implemented optionally to maintain object identity across frames.

- **Recommendation: DeepSORT:** We recommend this appearance-based tracker for its robustness in re-identifying vehicles after temporary occlusion. It works seamlessly with YOLO bounding boxes and maintains stable IDs with only a minimal FPS drop ($\approx 5\text{--}10\%$).
- **Alternative: ByteTrack:** A lightweight, motion-only tracker best suited for extreme high-FPS scenarios, though less reliable under long occlusions.

DeepSORT Implementation Snippet:

```
from deep_sort_realtime.deepsort_tracker import DeepSort
```

```
tracker = DeepSort(max_age=30, n_init=3, nn_budget=100)
tracks = tracker.update_tracks(detections, frame=frame)
```

4. Optimization Summary

The system's high performance is due to several key optimizations:

- **Dataset Expansion:** The dataset was significantly increased from 370 to 1,110 images using augmentations (motion blur, cutout) to improve robustness, particularly against blur and occlusion.
- **Model Architecture:** The chosen YOLOv10 variant ($\approx 11.1\text{M}$ parameters) and its multi-scale feature extraction effectively balance accuracy and speed across all vehicle sizes.
- **Training:** Optimal training was achieved at **100 epochs** (best mAP with no overfitting). The deployed model is named **yolov10.pt**.
- **Hardware/Computation:** Inference on dual NVIDIA T4 GPUs delivers a total speed of **$\approx 93\text{ FPS}$** , calculated from a per-image processing time of $\approx 10.7\text{ ms}$ ($\approx 7.7\text{ ms}$ for YOLOv10 + $\approx 3\text{ ms}$ for pre/post-processing).

5. Performance Analysis

- **Total Processing Time:** $\approx 10.7\text{ ms}$ per image.
- **Frame Rate:** **$\approx 93\text{ FPS}$** .
- **Conclusion:** This capability is fully real-time and ideal for live traffic monitoring.

5.2 Detection Accuracy (Test Set)

Metric	Value
Precision	0.732
Recall	0.667
mAP@50	0.752
mAP@50–95	0.552

- **Highlights:** We observed high mAP@50–95 for large vehicles (Bus: ≈0.615 / Truck: ≈0.603) and excellent precision for small objects (Motorcycle: ≈0.858).

5.3 Resource Efficiency

- **Model Size:** ≈11.13M parameters, ≈28.5 GFLOPs.
- **GPU Usage:** The system demonstrates efficient VRAM utilization on dual T4s, suitable for a moderate-resource real-time deployment.

6. Setup Instructions

6.1 Core Dependencies

```
ultralytics>=8.0.0
opencv-python>=4.8.0
pillow>=10.0.0
torch>=2.0.0
torchvision>=0.15.0
numpy>=1.24.0
streamlit>=1.28.0
Deep-sort-realtime
```

6.2 Installation Steps

1. **Clone & Navigate:** `git clone https://github.com/arnab9961/vehicle_detection.git; cd vehicle_detection.`
2. **Environment:** Create and activate a virtual environment: `python -m venv venv; source venv/bin/activate.`
3. **Install:** Install all required dependencies: `pip install -r requirements.txt`.
4. **Model Weights:** Ensure `yolov10.pt` (or current weights file) is placed in the project root.
5. **Run:** Launch the application: `streamlit run app.py`.

7. Application Features (Streamlit GUI)

The user interface provides comprehensive control and visualization:

- **Image Detection:** Drag-and-drop image upload, real-time bounding box display, and downloadable results.
- **Video Detection:** Custom video upload, frame-by-frame updates, annotated video output, and aggregated statistics.
- **Configuration:** An adjustable panel for confidence and IoU thresholds (0.0–1.0), and model path management.
- **Statistics Dashboard:** Displays total detections, class-wise counts, and performance metrics.

8. Model Training Details

8.2 Training Configuration (`data.yaml`)

```
train: ./dataset/images/train
val: ./dataset/images/val
test: ./dataset/images/test
nc: 11
names: [
    'Auto Rickshaw', 'Cycle Rickshaw', 'CNG / Tempo', 'Bus',
    'Jeep / SUV', 'Microbus', 'Minibus', 'Motorcycle',
    'Truck', 'Private Sedan Car', 'Trailer'
]
```

8.3 Training Command

```
from ultralytics import YOLO
model = YOLO('yolov10m.pt')
results = model.train(
    data='data.yaml',
    epochs=100,
    imgsz=640,
    batch=16,
    device=0, # GPU
    patience=20,
    project='vehicle_detection',
    name='yolov10_vehicles'
)
```

8.4 Key Hyperparameters

Parameter	Value
Learning rate	0.01
Momentum	0.937
Weight decay	0.0005

Batch size	16
Image size	640
Epochs	100

9. Future Enhancements

Our plan for further development includes:

- **Tracking Integration:** Implementing the full DeepSORT/ByteTrack pipeline with trajectory visualization.
- **Advanced Analytics:** Adding speed estimation, traffic flow analysis, and heatmaps.
- **Multi-camera Support:** Developing cross-camera tracking capabilities and a centralized dashboard.
- **Edge Deployment:** Optimization and documentation for low-resource platforms like NVIDIA Jetson and Raspberry Pi.

10. Conclusion

The final YOLOv10-based vehicle detection system successfully delivers on the goals of real-time performance (≈ 93 FPS) and strong accuracy (mAP@50 ≈ 0.752). Its robust performance and efficient design confirm its status as a production-ready solution for complex intelligent transportation challenges.